

Sowing Success: How Machine Learning Helps Farmers Select the Best Crops



Measuring essential soil metrics such as nitrogen, phosphorous, potassium levels, and pH value is an important aspect of assessing soil condition. However, it can be an expensive and time-consuming process, which can cause farmers to prioritize which metrics to measure based on their budget constraints.

Farmers have various options when it comes to deciding which crop to plant each season. Their primary objective is to maximize the yield of their crops, taking into account different factors. One crucial factor that affects crop growth is the condition of the soil in the field, which can be assessed by measuring basic elements such as nitrogen and potassium levels. Each crop has an ideal soil condition that ensures optimal growth and maximum yield.

A farmer reached out to you as a machine learning expert for assistance in selecting the best crop for his field. They've provided you with a dataset called `soil_measures.csv`, which contains:

- `"N"` : Nitrogen content ratio in the soil
- `"P"` : Phosphorous content ratio in the soil
- `"K"` : Potassium content ratio in the soil
- `"pH"` value of the soil
- `"crop"` : categorical values that contain various crops (target variable).

Each row in this dataset represents various measures of the soil in a particular field. Based on these measurements, the crop specified in the `"crop"` column is the optimal choice for that field.

In this project, you will build multi-class classification models to predict the type of `"crop"` and identify the single most importance feature for predictive performance.

```
# All required libraries are imported here for you.
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics

# Load the dataset
crops = pd.read_csv("soil_measures.csv")

# Write your code here
crops.head()
```

index	...	↑↓	N	...	↑↓	P	...	↑↓	K	...	↑↓	ph	...	↑↓
		0			90			42			43			6.502985292
		1			85			58			41			7.038096361
		2			60			55			44			7.840207144
		3			74			35			40			6.980400905
		4			78			42			42			7.628472891

Rows: 5

[Expand](#)

```
crops.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2200 entries, 0 to 2199
Data columns (total 5 columns):
Column Non-Null Count Dtype
--- ---
0 N 2200 non-null int64
1 P 2200 non-null int64
2 K 2200 non-null int64
3 ph 2200 non-null float64
4 crop 2200 non-null object
dtypes: float64(1), int64(3), object(1)
memory usage: 86.1+ KB

crops.describe()

index	...	↑↓	N	...	↑↓	P	...	↑↓	K	...	↑↓	ph
count			2200			2200			2200			
mean			50.5518181818			53.3627272727			48.1490909091			
std			36.9173338338			32.9858827386			50.6479305467			
min			0			5			5			
25%			21			28			20			
50%			37			51			32			
75%			84.25			68			49			
max			140			145			205			

Rows: 8 Expand

🔗 crops.isna.sum

```
# Check for missing values in the crops DataFrame  
crops.isna().sum()  
crops["crop"].value_counts()
```

ind...	...	↑↓	...	↑↓
rice			100	
maize			100	
jute			100	
cotton			100	
coconut			100	
papaya			100	
orange			100	
apple			100	
muskmelon			100	
watermelon			100	
grapes			100	
mango			100	
banana			100	
pomegranate			100	
lentil			100	
blackgram			100	

Rows: 22 Expand

```
#splitting the data  
y= crops['crop']
```

How likely are you to recommend DataLab to a friend or co-worker?

Not at all likely

0

1

2

3

4

5

6

7

8

9

10

Extremely likely

```
# Create a dictionary to store the model performance for each feature
feature_performance = {}
```

```
# Fix: Use the correct column names as present in X_train/X_test
print("X_train columns:", X_train.columns) # Debug: See actual column names
```

```
# Map the feature names to the actual column names in X_train/X_test
```

```
feature_map = {
    "N": "N",      # update if actual column name is different
    "P": "P",      # update if actual column name is different
    "K": "K",      # update if actual column name is different
    "ph": "ph"     # update if actual column name is different
}
```

```
feature_performance = {}
```

```
for feature in ["N", "P", "K", "ph"]:
    col_name = feature_map[feature]
    if col_name not in X_train.columns:
        print(f"Column '{col_name}' not found in X_train. Skipping.")
        continue
    log_reg = LogisticRegression(multi_class='multinomial')
    log_reg.fit(X_train[[col_name]], y_train)
    y_pred = log_reg.predict(X_test[[col_name]])
    f1 = metrics.f1_score(y_test, y_pred, average='weighted')
    feature_performance[feature] = f1
    print(f"F1.score for {feature}: {f1}")
```

```
X_train columns: Index(['N', 'P', 'K', 'ph'], dtype='object')
F1.score for N: 0.09149868209906838
F1.score for P: 0.14761942909728204
F1.score for K: 0.23896974566001802
F1.score for ph: 0.04532731061152114
```

```
best_predictive_feature = {"K": feature_performance['K']}
best_predictive_feature
```

```
{'K': 0.23896974566001802}
```

```
# getting dummies only for crop column
crops_dummy = pd.get_dummies(crops['crop'], drop_first=True)
crops_dummy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2200 entries, 0 to 2199
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	banana	2200 non-null	uint8
1	blackgram	2200 non-null	uint8
2	chickpea	2200 non-null	uint8
3	coconut	2200 non-null	uint8
4	coffee	2200 non-null	uint8
5	cotton	2200 non-null	uint8
6	grapes	2200 non-null	uint8
7	jute	2200 non-null	uint8
8	kidneybeans	2200 non-null	uint8
9	lentil	2200 non-null	uint8
10	maize	2200 non-null	uint8
11	mango	2200 non-null	uint8
12	mothbeans	2200 non-null	uint8
13	mungbean	2200 non-null	uint8
14	muskmelon	2200 non-null	uint8
15	orange	2200 non-null	uint8
16	pomegranate	2200 non-null	uint8

How likely are you to recommend DataLab to a friend or co-worker?

Not at all likely 0 1 2 3 4 5 6 7 8 9 10 Extremely likely

```
#crops_dummy = pd.concat([crops, crops_dummy], axis=1)
#crops_dummy = crops_dummy.drop("crop", axis=1)
#crops_dummy
```

```
crops_dummy = pd.concat([crops,crops_dummy],axis =1)
crops_dummy = crops_dummy.drop('crop',axis =1)
crops_dummy
```

...	↑↓	...	↑↓	...	↑↓	...	↑↓	ph	...	↑↓	...	↑↓	b.	...	↑↓	...	↑↓	...	↑↓	...	↑↓	...	↑↓	...	↑↓	...	↑↓	...	↑↓	kid...
0		90		42		43		6.502985292		0		0		0		0		0		0		0		0		0		0		
1		85		58		41		7.038096361		0		0		0		0		0		0		0		0		0		0		
2		60		55		44		7.840207144		0		0		0		0		0		0		0		0		0		0		
3		74		35		40		6.980400905		0		0		0		0		0		0		0		0		0		0		
4		78		42		42		7.628472891		0		0		0		0		0		0		0		0		0		0		
5		69		37		42		7.073453503		0		0		0		0		0		0		0		0		0		0		
6		69		55		38		5.70080568		0		0		0		0		0		0		0		0		0		0		
7		94		53		40		5.718627178		0		0		0		0		0		0		0		0		0		0		
8		89		54		38		6.685346424		0		0		0		0		0		0		0		0		0		0		
9		68		58		38		6.336253525		0		0		0		0		0		0		0		0		0		0		
10		91		53		40		5.386167788		0		0		0		0		0		0		0		0		0		0		
11		90		46		42		7.50283396		0		0		0		0		0		0		0		0		0		0		
12		78		58		44		5.108681786		0		0		0		0		0		0		0		0		0		0		
13		93		56		36		6.98435366		0		0		0		0		0		0		0		0		0		0		
14		94		50		37		6.94801983		0		0		0		0		0		0		0		0		0		0		
15		60		48		39		7.042299069		0		0		0		0		0		0		0		0		0		0		

Rows: 2,200

Expand

```
final_feature = ['N', 'K', 'ph']
X_train, X_test, y_train, y_test = train_test_split(
    crops[final_features],
    crops["crop"],
    test_size=0.2,
    random_state=42
)
log_reg = LogisticRegression(max_iter=2000, multi_class='multinomial')
log_reg.fit(X_train, y_train)
y_pred = log_reg.predict(X_test)
f_error = f1_score(y_test, y_pred, average='weighted')
model_performance = f1_score(y_test, y_pred, average="weighted")
print(model_performance)
```

0.558010495235685

How likely are you to recommend DataLab to a friend or co-worker?

X

Not at all likely

0

1

2

3

4

5

6

7

8

9

10

Extremely likely